

An Introduction to Minimum Principle

Nuthasith (Nat) Gerdpratoom

2024/07/24

Control & Optimization Laboratory

Table of contents

1. Minimum Principle and Lagrange Multipliers
2. The Penalty Approach
3. Application to LQR
4. Nonlinear Examples

Minimum Principle and Lagrange Multipliers

Lagrange Multipliers in Finite Dimensions

Consider an optimization problem with constraint $g(x) = \mathbf{0}$, where $g : \mathbb{R}^m \rightarrow \mathbb{R}^d$:

$$\begin{aligned} \min_x V(x) \\ \text{s.t. } g(x) = \mathbf{0} \end{aligned}$$

In constrained optimization in \mathbb{R}^m , we use Lagrange multipliers:

$$\hat{V}(x, p) = V(x) + p^T g(x)$$

where p is a Lagrange multiplier. If x^0, p^0 is a stationary point, then:

$$\nabla_x \hat{V}(x^0, p^0) = \nabla_x V(x^0) + \nabla_x g(x^0) p^0 = \mathbf{0}$$

$$\nabla_p \hat{V}(x^0, p^0) = g(x^0) = \mathbf{0}$$

Lagrange Multipliers in Finite Dimensions




Figure 1: An optimization problem in \mathbb{R}^2 with a single constraint, $g(x) = 0$.

$$\nabla_x \hat{V}(x^0, p^0) = \nabla_x V(x^0) + \nabla_x g(x^0) p^0 = \mathbf{0}$$

$$\nabla_p \hat{V}(x^0, p^0) = g(x^0) = \mathbf{0}$$

Generalizing to Infinite Dimensions

To generalize this to functional minimization, suppose that F is a functional on $D^r[t_0, t_1]$, $F(z)$ is a real number, and $z \in D^r[t_0, t_1]$. We define the directional derivative:

$$D_\eta F(z) = \lim_{\epsilon \rightarrow 0} \frac{F(z + \epsilon \eta) - F(z)}{\epsilon}$$

A function z_0 is a stationary point of F if $D_\eta F(z_0) = 0$ for any η .




Figure 2: A Perturbation of function $z \in D[t_0, t_1]$

Steps to Extend the Lagrange Multiplier Method

Let us consider a constrained functional optimization problem:

$$\begin{aligned} \min_u V(x, u) &= \int \ell d\tau + m \\ \text{s.t. } \dot{x} - f &= \mathbf{0}, \quad x \in D^n[t_0, t_1], u \in D^m[t_0, t_1]. \end{aligned}$$

1. **Append state equations:** Define a new cost functional including the state dynamics:

$$\hat{V}(x, u) = \int_{t_0}^{t_1} \ell dt + m(x(t_1)) + \int_{t_0}^{t_1} p^T (f - \dot{x}) dt$$

2. **Integration by parts:** Eliminate the derivative of x :

$$\int_{t_0}^{t_1} p^T \dot{x} dt = p^T x|_{t_0}^{t_1} - \int_{t_0}^{t_1} \dot{p}^T x dt$$

Steps to Extend the Lagrange Multiplier Method (cont.)

3. Define the Hamiltonian:

$$H(x, p, u, t) := \ell(x, u, t) + p^T f(x, u, t)$$

$$\begin{aligned}\hat{V}(x, u) = & \int_{t_0}^{t_1} H(x, p, u, t) dt + \int_{t_0}^{t_1} \dot{p}^T x dt \\ & + p^T(t_0)x(t_0) - p^T(t_1)x(t_1) + m(x(t_1))\end{aligned}$$

4. **Compute variations:** For perturbations $x(t, \epsilon) = x^\circ(t) + \epsilon\eta(t)$ and $u(t, \delta) = u^\circ(t) + \delta\psi(t)$, let $\hat{V}(\epsilon) = \hat{V}(x^\circ + \epsilon\eta, u^\circ)$, and $\hat{V}(\delta) = \hat{V}(x^\circ, u^\circ + \delta\psi)$.

$$\begin{aligned}\frac{d}{d\epsilon} \hat{V}(\epsilon) = & \int_{t_0}^{t_1} \frac{\partial}{\partial x} H(x^\circ, p, u^\circ, t) \eta(t) dt + \int_{t_0}^{t_1} \dot{p}^T \eta(t) dt \\ & - p^T(t_1)\eta(t_1) + p^T(t_0)\eta(t_0) + \frac{\partial}{\partial x} m(x^\circ(t_1))\eta(t_1) = 0.\end{aligned}$$

Similarly, by considering perturbations in u° :

$$\frac{d}{d\delta} \hat{V}(\delta) = \int_{t_0}^{t_1} \frac{\partial}{\partial u} H(x^\circ, p, u^\circ, t) \psi(t) dt = 0$$

Steps to Extend the Lagrange Multiplier Method (cont.)

4. **Compute variations:** Determine perturbations in x° :

$$\int_{t_0}^{t_1} \frac{\partial}{\partial x} H(x^\circ, p, u^\circ, t) \eta(t) dt + \int_{t_0}^{t_1} \dot{p}^T \eta(t) dt \\ - p^T(t_1) \eta(t_1) + p^T(t_0) \eta(t_0) + \frac{\partial}{\partial x} m(x^\circ(t_1)) \eta(t_1) = 0.$$

Similarly, by considering perturbations in u° :

$$\int_{t_0}^{t_1} \frac{\partial}{\partial u} H(x^\circ, p, u^\circ, t) \psi(t) dt = 0$$

5. **Derive necessary conditions:** From the variations, derive the equations:

$$\dot{p} = -\nabla_x H, \quad p(t_1) = \nabla_x m(x(t_1))$$

$$\nabla_u H = 0$$

Deriving the Minimum Principle

Combine the steps to derive the Minimum Principle. If u° is optimal, there exists a costate $p(t)$ such that:

$$u^\circ(t) \in \arg \min_u H(x^\circ(t), p(t), u, t)$$

The state and costate satisfy:

$$\dot{x}^\circ = \nabla_p H, \quad \dot{p} = -\nabla_x H$$

with boundary conditions:

$$x(t_0) = x_0, \quad p(t_1) = \nabla_x m(x(t_1))$$

The Penalty Approach

The Penalty Approach

Another approach to the Minimum Principle involves relaxing the hard constraint $\dot{x} - f = 0$, and instead imposing a large, yet "soft" constraint by defining the cost function:

$$\hat{V}(x, u) = \int_{t_0}^{t_1} \ell(x(t), u(t), t) dt + \frac{k}{2} \int_{t_0}^{t_1} |\dot{x}(t) - f(x(t), u(t), t)|^2 dt + m(x(t_1))$$

If (x_k, u_k) minimizes \hat{V}_k , and letting (x°, u°) denote a solution to the original problem:

$$\hat{V}_k(x_k, u_k) \leq \hat{V}_k(x^\circ, u^\circ) = V^\circ$$

Assuming ℓ and m are positive, subtracting the left side of the above inequality gives the uniform bound:

$$\int_{t_0}^{t_1} |\dot{x}(t) - f(x(t), u(t), t)|^2 dt \leq \frac{2}{k} V^\circ$$

Implications for Large k

As k becomes large, the term $\frac{k}{2} \int_{t_0}^{t_1} |\dot{x}(t) - f(x(t), u(t), t)|^2 dt$ ensures that $\dot{x}(t) \approx f(x(t), u(t), t)$:

$$\int_{t_0}^{t_1} |\dot{x}(t) - f(x(t), u(t), t)|^2 dt \rightarrow 0 \text{ as } k \rightarrow \infty$$

Thus, for large k , the pair (x_k, u_k) will approximately satisfy the differential equation $\dot{x} = f$.

If we perturb x_k to form $x_k + \epsilon\eta$ and define $\hat{V}(\epsilon) = \hat{V}(x_k + \epsilon\eta, u_k)$, then we must have:

$$\left. \frac{d}{d\epsilon} \hat{V}(\epsilon) \right|_{\epsilon=0} = 0$$

Using the definition of \hat{V} , we get:

$$\begin{aligned} \hat{V}(\epsilon) = & \int_{t_0}^{t_1} \ell(x_k(t) + \epsilon\eta(t), u_k(t), t) dt + m(x_k(t_1) + \epsilon\eta(t_1)) \\ & + \frac{k}{2} \int_{t_0}^{t_1} |\dot{x}_k(t) + \epsilon\dot{\eta}(t) - f(x_k(t) + \epsilon\eta(t), u_k(t), t)|^2 dt \end{aligned}$$

Computing the Derivative

The derivative of this expression with respect to ϵ can be computed as follows:

$$\begin{aligned}\frac{d}{d\epsilon} \hat{V}(0) &= \int_{t_0}^{t_1} \frac{\partial \ell}{\partial x}(x_k(t), u_k(t), t) \eta(t) dt \\ &\quad + k \int_{t_0}^{t_1} (\dot{x}_k(t) - f(x_k(t), u_k(t), t))^T \left[\dot{\eta}(t) - \frac{\partial f}{\partial x}(x_k(t), u_k(t), t) \eta(t) \right] dt \\ &\quad + \frac{\partial m}{\partial x}(x_k(t_1)) \eta(t_1) \\ &= 0\end{aligned}$$

Computing the Derivative (cont.)

To eliminate the derivative term $\dot{\eta}$, we integrate by parts:

$$\begin{aligned} & \int_{t_0}^{t_1} \left\{ \frac{\partial \ell}{\partial x}(x_k(t), u_k(t), t) + p_k(t)^T \frac{\partial f}{\partial x}(x_k(t), u_k(t), t) + \frac{d}{dt} (p_k(t)^T) \right\} \eta(t) dt \\ & - p_k(t_1)^T \eta(t_1) + \frac{\partial m}{\partial x}(x_k(t_1)) \eta(t_1) \\ & = 0 \end{aligned}$$

where we have set $p_k(t) = -k(\dot{x}_k(t) - f(x_k(t), u_k(t), t))$.

Resulting Equations

Since η is arbitrary, we have:

$$\mathbf{0} = \frac{d}{dt} (p_k(t)^T) + \frac{\partial \ell}{\partial x}(x_k(t), u_k(t), t) + p_k(t)^T \frac{\partial f}{\partial x}(x_k(t), u_k(t), t)$$

$$\mathbf{0} = \frac{d}{dt} (p_k(t)^T) + \frac{\partial H}{\partial x}(x_k(t), p_k(t), u_k(t), t)$$

$$\dot{p}_k = -\nabla_x H$$

with the boundary condition:

$$p_k(t_1) = \nabla_x m(x_k(t_1))$$

Considering perturbations in u gives the equation:

$$\nabla_u H(x_k(t), p_k(t), u_k(t), t) = \mathbf{0}$$

This is a weak form of the Minimum Principle for the perturbed problem.

Application to LQR

The LQR problem tests the Minimum Principle's utility for constructing optimal policies. Consider the general LTI model with quadratic cost:

$$\dot{x} = Ax + Bu$$

$$V = \int_{t_0}^{t_1} (x^T Q x + u^T R u) dt + x^T(t_1) M x(t_1)$$

Solving the LQR Problem

The Hamiltonian:

$$H = x^T Q x + u^T R u + p^T (A x + B u)$$

Control can be computed through:

$$\nabla_u H = 0 \Rightarrow u = -\frac{1}{2} R^{-1} B^T p$$

Resulting in:

$$\dot{x} = A x + B u = A x - \frac{1}{2} B R^{-1} B^T p$$

Through the expression $\nabla_x H = 2Qx + A^T p$, we find that \dot{p} is:

$$\dot{p} = -\nabla_x H = -2Qx - A^T p$$

Solving the LQR Problem (cont.)

The equations form the coupled set of differential equations:

$$\begin{pmatrix} \dot{x} \\ \dot{p} \end{pmatrix} = \begin{pmatrix} A & -\frac{1}{2}BR^{-1}B^T \\ -2Q & -A^T \end{pmatrix} \begin{pmatrix} x \\ p \end{pmatrix}$$

with boundary conditions:

$$x(t_0) = x_0$$

$$p(t_1) = \nabla_x m(x_k(t_1)) = 2Mx(t_1)$$

If we scale p by defining $\lambda = \frac{1}{2}p$, we get:

$$\begin{pmatrix} \dot{x} \\ \dot{\lambda} \end{pmatrix} = \begin{pmatrix} A & -BR^{-1}B^T \\ -Q & -A^T \end{pmatrix} \begin{pmatrix} x \\ \lambda \end{pmatrix} = \mathcal{H} \begin{pmatrix} x \\ \lambda \end{pmatrix}$$

with the optimal control:

$$u^o(t) = -R^{-1}B^T\lambda(t)$$

Solving the ODE

This ODE can be solved using the *sweep method*. Suppose that $\lambda(t) = P(t)x(t)$. Then:

$$\dot{\lambda} = \dot{P}x + P(Ax + Bu)$$

Substituting $u^o = -R^{-1}B^T\lambda = -R^{-1}B^TPx$ gives:

$$\dot{\lambda} = \dot{P}x + P(Ax - BR^{-1}B^TPx)$$

From the coupled ODE with \mathcal{H} above, we also have:

$$\dot{\lambda} = -Qx - A^T\lambda = -Qx - A^TPx$$

Equating the two expressions for $\dot{\lambda}$ gives the Riccati Differential Equation:

$$\dot{P}x + P(Ax - BR^{-1}B^TPx) = -Qx - A^TPx$$

The boundary condition for λ is:

$$\lambda = \frac{1}{2}p = \frac{1}{2}\nabla_x m(x(t_1)) \implies P(t_1) = M$$

Solving for $P(t)$ gives $\lambda(t)$, which in turn gives $p(t)$.

Nonlinear Examples

Minimum Principle with Input Constraints

Theorem 11.2. Suppose that $x(t_1)$ is free, t_1 is fixed, and suppose that u° is a solution to the optimal control problem:

$$V^\circ = \min_{u \in \mathcal{U}} V(u)$$

We then have

(a) There exists a costate vector $p(t)$ such that

$$u^\circ(t) = \arg \min_{u \in \mathcal{U}} H(x^\circ(t), p(t), u, t).$$

(b) The pair (p, x°) satisfy the 2-point boundary value problem:

$$\begin{aligned}\dot{x}^\circ(t) &= \nabla_p H(x^\circ(t), p(t), u^\circ(t), t) \quad (= f(x^\circ(t), u^\circ(t), t)) \\ \dot{p}(t) &= -\nabla_x H(x^\circ(t), p(t), u^\circ(t), t),\end{aligned}$$

with the two boundary conditions

$$x(t_0) = x_0; \quad p(t_1) = \nabla_x m(x(t_1), t).$$

Example 11.5.1: Bilinear System

Consider the control of a bilinear system, defined by the differential equation:

$$\dot{x} = ux, \quad x(0) = x_0 > 0, \quad 0 \leq u(t) \leq 1$$

Suppose the goal is to make x large while keeping the derivative of x small on average. The cost criterion is:

$$V(u) = \int_0^{t_1} \dot{x}(\tau) - x(\tau) d\tau = \int_0^{t_1} [u(\tau) - 1]x(\tau) d\tau$$

The Hamiltonian becomes:

$$H(x, p, u, t) = pf + \ell = x\{u(p + 1) - 1\}$$

Example 11.5.1: Bilinear System (cont.)

By the Minimum Principle:

$$\begin{aligned}H(x, p, u, t) &= x\{u(p+1) - 1\} \\ \nabla_u H &= x(p+1)\end{aligned}$$

Since $x^\circ(t) > 0$ for all t , the minimization leads to:

$$u^\circ(t) = \begin{cases} 1 & \text{if } p(t) + 1 < 0 \\ 0 & \text{if } p(t) + 1 > 0 \\ \text{unknown} & \text{if } p(t) + 1 = 0 \end{cases}$$

Example 11.5.1: Bilinear System (cont.)

The costate and state variables satisfy the differential equations:

$$\begin{aligned}\dot{p} &= -\frac{\partial H}{\partial x} = -(p+1)u^\circ + 1 \\ \dot{x} &= \frac{\partial H}{\partial p} = u^\circ x^\circ\end{aligned}$$

with boundary conditions:

$$x^\circ(0) = x_0, \quad p(t_1) = \frac{\partial m}{\partial x} = 0$$

Example 11.5.1: Bilinear System (cont.)

Solving the Differential Equations

If $t = t_1$, then $p(t) + 1 = 1 > 0$ so $u^\circ(t) = 0$. By continuity, if $t \approx t_1$, then $p(t) > 0$, so $u^\circ(t) = 0$ and:

$$\dot{p}(t) = 1 \quad \text{for } t \approx t_1 \implies p(t) = t - t_1$$

For $t < t_1 - 1$, we have $p(t) + 1 < 0$ so $u^\circ(t) = 1$. Since $p(t_1 - 1) = -1$, we have:

$$\dot{p} = -p \quad \text{for } t < t_1 - 1 \implies p(t) = -e^{-t+t_1-1}$$

Example 11.5.1: Bilinear System (cont.)




Figure 3: The costate trajectory, optimal control, and optimal state trajectory for the bilinear model.

The optimal state trajectory is:

$$x^o(t) = \begin{cases} x_0 e^t & \text{if } t < t_1 - 1 \\ x_0 e^{t_1 - 1} & \text{if } t > t_1 - 1 \end{cases}$$

Minimum Principle with Final Value Constraints

Theorem 11.3. Suppose that $t_0, t_1, x(t_0) = x_0$, and $x(t_1) = x_1$ are prespecified, and suppose that u° is a solution to the optimal control

$$\dot{x}(t) = f(x(t), u(t), t), \quad x(t_0) = x_0 \in \mathbb{R}^n$$

$$V(u) = \int_{t_0}^{t_1} \ell(x(t), u(t), t) dt + m(x(t_1)),$$

subject to those constraints. Then

- (a) There exists a costate vector $p(t)$ such that

$$u^\circ(t) = \arg \min_u H(x^\circ(t), p(t), u, t).$$

- (b) The pair (p, x°) satisfy the 2-point boundary value problem

$$\dot{x}^\circ = \nabla_p H, \quad \dot{p} = -\nabla_x H,$$

with the two boundary conditions

$$x(t_0) = x_0; \quad x(t_1) = x_1.$$

Example 11.5.2: LQR Problem with Terminal State

Consider the LQR problem with the terminal state specified, $x(t_1) = x_1$.
The cost criterion is:

$$V = \frac{1}{2} \int_0^{t_1} (x^T Q x + u^T R u) dt, \quad R \succ 0, Q \succeq 0$$

The Hamiltonian:

$$H = p(Ax + Bu) + \frac{1}{2}(x^T Q x + u^T R u)$$

The optimal control:

$$u^o(t) = -R^{-1}B^T p(t)$$

Example 11.5.2: LQR Problem with Terminal State (cont.)

The costate vector satisfies the differential equation:

$$\dot{p}(t) = -\frac{\partial H}{\partial x} = -p^T A - x^T Q$$

$$\begin{pmatrix} \dot{x} \\ \dot{p} \end{pmatrix} = \begin{pmatrix} A & -BR^{-1}B^T \\ -Q & -A^T \end{pmatrix} \begin{pmatrix} x \\ p \end{pmatrix}$$

with boundary conditions:

$$x(t_0) = x_0, \quad x(t_1) = x_1$$

Example 11.5.2: LQR Problem with Terminal State (cont.)

Let $\psi(t, \tau)$ denote the state transition matrix:

$$\frac{d}{dt}\psi(t, \tau) = \mathcal{H}(t)\psi(t, \tau), \quad \psi(t, t) = I$$

Decompose $\psi(t, \tau)$ as:

$$\psi(t, \tau) = \begin{pmatrix} \psi_{11}(t, \tau) & \psi_{12}(t, \tau) \\ \psi_{21}(t, \tau) & \psi_{22}(t, \tau) \end{pmatrix}$$

Example 11.5.2: LQR Problem with Terminal State (cont.)

Compute the initial condition $p(t_0) = p_0$ to find $p(t)$ for all t :

$$x_1 = x(t_1) = \psi_{11}(t_1, t_0)x_0 + \psi_{12}(t_1, t_0)p_0$$

Assuming $\psi_{12}(t_1, t_0)$ is invertible, we have:

$$p_0 = \psi_{12}(t_1, t_0)^{-1}(x_1 - \psi_{11}(t_1, t_0)x_0)$$

Example 11.5.2: LQR Problem with Terminal State (cont.)

For all t :

$$p(t) = \psi_{21}(t, t_0)x_0 + \psi_{22}(t, t_0)p_0$$

The optimal control is:

$$u^\circ(t) = -R^{-1}B^T p(t)$$

The optimal state trajectory is:

$$x^\circ(t) = \psi_{11}(t, t_0)x_0 + \psi_{12}(t, t_0)p_0, \quad t \geq t_0$$

Suppose $t_0, x(t_0) = x_0$ are fixed, and some components of $x(t_1)$ are specified. The Minimum Principle with free terminal time states:

1. There exists a costate vector $p(t)$ such that
$$u^\circ(t) = \arg \min_u H(x^\circ(t), p(t), u, t).$$
2. The pair (p, x°) satisfy the 2-point boundary value problem with modified boundary conditions.
3. The unspecified terminal time t_1 satisfies:

$$\frac{\partial m}{\partial t}(x^\circ(t_1), t) + H(x^\circ(t_1), p(t_1), u^\circ(t_1), t_1) = 0$$

Example 11.5.3: Minimum Time Problem

Consider a single input linear state space model:

$$\dot{x} = Ax + bu$$

We wish to find u° which drives x from $x(0) = x_0$ to $x(t_1) = x_1$ in minimum time, under the constraint $|u(t)| \leq 1$. The cost criterion is:

$$V(u) = t_1 = \int_0^{t_1} 1 dt$$

The Hamiltonian is:

$$H = 1 + p^T(Ax + bu)$$

Minimizing H gives a bang-bang control law:

$$u^{\circ}(t) = \begin{cases} 1 & \text{if } p(t)^T b < 0 \\ -1 & \text{if } p(t)^T b > 0 \end{cases}$$

If A, b are time-invariant and the eigenvalues of A are real, distinct, and negative, then $p(t)^T b$ changes sign at most $n - 1$ times, bounding the number of switching times for the optimal control.

The costate equation is:

$$\dot{p} = -\nabla_x H = -A^T p$$

Using the final time boundary condition $H|_{t=t_1} = 0$:

$$1 + p(t_1)^T A x^\circ(t_1) + p(t_1)^T b u^\circ(t_1) = 1 + p(t_1)^T A x^\circ(t_1) - |p(t_1)^T b| = 0$$

Example: Double Integrator

Consider the double integrator $\ddot{y} = u$:

$$A = \begin{pmatrix} 0 & 1 \\ 0 & 0 \end{pmatrix}, \quad b = \begin{pmatrix} 0 \\ 1 \end{pmatrix}$$

The optimal control is:

$$u^\circ(t) = -\text{sgn}(p(t)^T b) = -\text{sgn}(p_2(t))$$

From the costate equations, for constants c_1, c_2 :

$$p_1(t) = c_1, \quad p_2(t) = -c_1 t + c_2$$

The optimal control is bang-bang:

$$u^\circ(t) = \begin{cases} 1 & \text{if } t < t_1 - 1 \\ 0 & \text{if } t > t_1 - 1 \end{cases}$$

The optimal state trajectories follow quadratic paths:

$$x_1(t) = \begin{cases} \frac{1}{2}(x_2(t))^2 + K_1 & \text{if } u^\circ(t) = 1 \\ -\frac{1}{2}(x_2(t))^2 + K_2 & \text{if } u^\circ(t) = -1 \end{cases}$$